

~13  
(↓)

# Computer Networks II

## Application Layer

Web and Electronic Mail

Amitangshu Pal  
Computer Science and Engineering  
IIT Kanpur


# Web and HTTP

---

# Web and HTTP

- **Web page** consists of **objects**
- Object can be HTML file, JPEG image, Java applet, audio file,...
- Web page consists of **base HTML-file** which includes **several referenced objects**
- Each object is addressable by a **URL**, e.g.,

<http://en.Wikipedia.org/somePage/pic.gif>



The diagram shows the URL <http://en.Wikipedia.org/somePage/pic.gif> with three curly braces underneath it. The first brace is under 'http' and labeled 'Protocol'. The second brace is under 'en.Wikipedia.org' and labeled 'Server'. The third brace is under '/somePage/pic.gif' and labeled 'Page on server'.

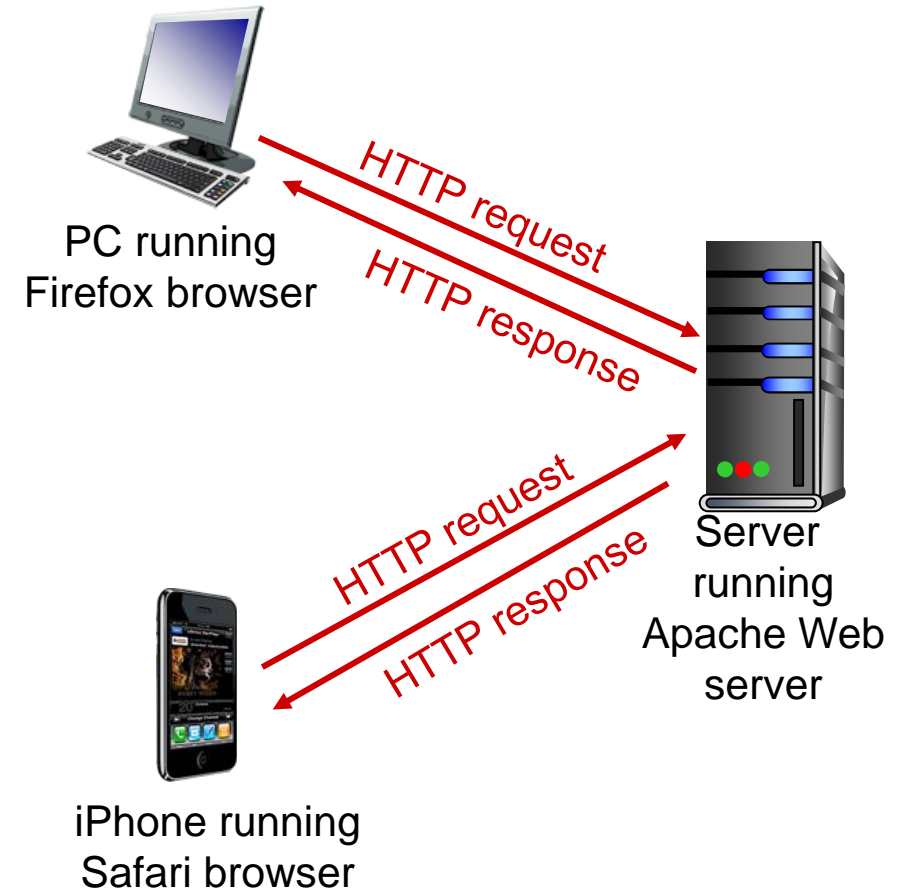
Protocol      Server      Page on server

---

# HTTP Overview

## HTTP: Hypertext transfer protocol

- Web's application layer protocol
- Client/server model
  - **Client:** Browser that requests, receives, (using HTTP protocol) and “displays” Web objects
  - **Server:** Web server sends (using HTTP protocol) objects in response to requests



# HTTP Overview

## Uses TCP:

- Client resolves server to IP address (DNS)
- Client initiates TCP connection (creates socket) to server, port 80
- Server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

## HTTP is “stateless”

- Server maintains no information about past client requests
-

# HTTP Request Message

- Two types of HTTP messages: **request, response**

Request line (GET, POST, HEAD commands) → GET /index.html HTTP/1.1\r\n

Header lines → Host: www-net.cs.umass.edu\r\n  
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10.15; rv:80.0) Gecko/20100101 Firefox/80.0 \r\n  
Accept: text/html,application/xhtml+xml\r\n  
Accept-Language: en-us,en;q=0.5\r\n  
Accept-Encoding: gzip,deflate\r\n  
Connection: keep-alive\r\n

Carriage return, line feed at start of line indicates end of header lines → \r\n

carriage return character  
line-feed character

The diagram illustrates the structure of an HTTP request message. It shows a request line followed by several header lines, each ending with a carriage return and line feed (\r\n). Annotations with arrows point to these characters: one points to the \r\n at the end of the request line, and another points to the \r\n at the end of the first header line. A third annotation points to the \r\n at the end of the last header line, stating that it indicates the end of the header lines.

# HTTP Response Message

Status line (protocol  
status code status phrase)

Header lines

Data, e.g., requested  
HTML file

```
HTTP/1.1 200 OK
Date: Tue, 08 Sep 2020 00:53:20 GMT
Server: Apache/2.4.6 (CentOS)
       OpenSSL/1.0.2k-fips PHP/7.4.9
       mod_perl/2.0.11 Perl/v5.16.3
Last-Modified: Tue, 01 Mar 2016 18:57:50 GMT
ETag: "a5b-52d015789ee9e"
Accept-Ranges: bytes
Content-Length: 2651
Content-Type: text/html; charset=UTF-8
\r\n
data data data data data ...
```

The diagram illustrates the structure of an HTTP response message. It consists of three main parts: the status line, header lines, and the data body. The status line is the first line of the message, followed by header lines, and then the data body. The status line is labeled 'Status line (protocol status code status phrase)' and points to 'HTTP/1.1 200 OK'. The header lines are labeled 'Header lines' and point to a block of lines including 'Date: Tue, 08 Sep 2020 00:53:20 GMT', 'Server: Apache/2.4.6 (CentOS)', 'OpenSSL/1.0.2k-fips PHP/7.4.9', 'mod\_perl/2.0.11 Perl/v5.16.3', 'Last-Modified: Tue, 01 Mar 2016 18:57:50 GMT', 'ETag: "a5b-52d015789ee9e"', 'Accept-Ranges: bytes', 'Content-Length: 2651', and 'Content-Type: text/html; charset=UTF-8'. The data body is labeled 'Data, e.g., requested HTML file' and points to the final line of the message, which is 'data data data data data ...'.

# HTTP Response Status Codes

- Status code appears in 1st line in server-to-client response message

- Some sample codes:

## 200 OK

- Request succeeded, requested object later in this message

## 301 Moved Permanently

- Requested object moved, new location specified later in this message (in Location: field)

## 400 Bad Request

- Request msg not understood by server

## 404 Not Found

- Requested document not found on this server

## 505 HTTP Version Not Supported


---



# Non-persistent HTTP (HTTP 1.0)

## Non-persistent HTTP


- At most one object sent over TCP connection
  - Connection then closed
- Downloading multiple objects required multiple connections



# Non-persistent HTTP (HTTP 1.0): Response Time

## HTTP response time:


- One RTT to initiate TCP connection
  - One RTT for HTTP request and first few bytes of HTTP response to return
  - File transmission time
- Non-persistent HTTP response time =  $2RTT + \text{file transmission time}$



# Non-persistent HTTP (HTTP 1.0): Response Time

## Non-persistent HTTP issues:

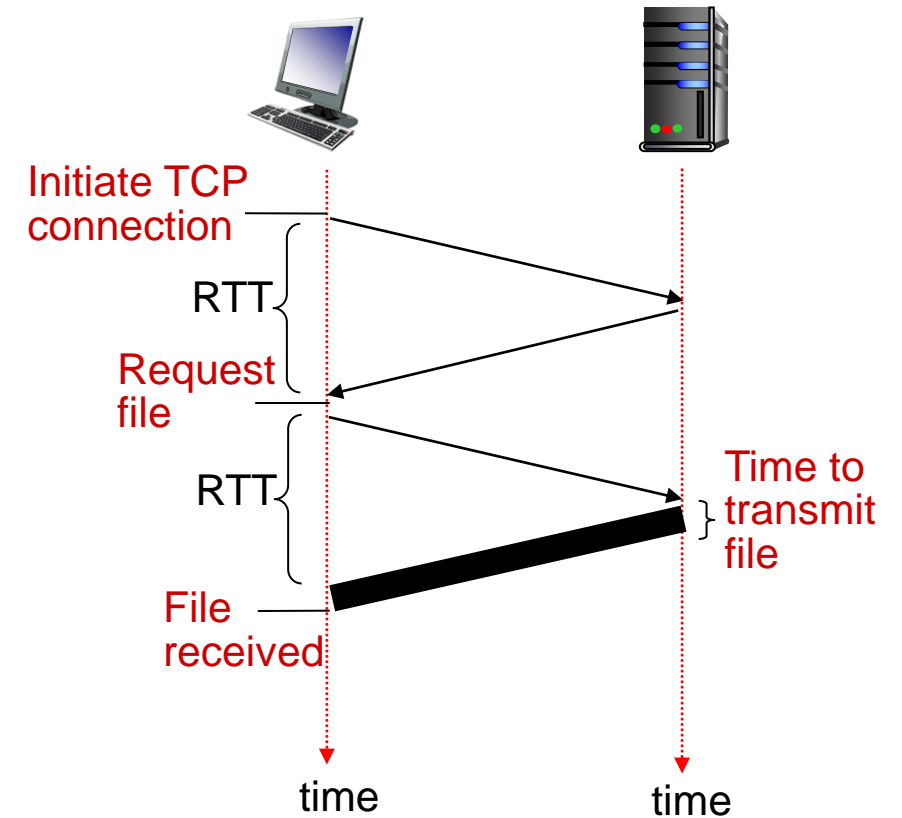
- Requires 2 RTTs per object
- OS overhead for each TCP connection
- Browsers often open parallel TCP connections to fetch referenced objects



# Non-persistent HTTP (HTTP 1.0): Response Time

We are downloading a webpage having a base HTML file with 5 embedded objects:


- All objects are from the same server
- The transmission time is negligible
- Objects are requested sequentially
- Object requests are made in parallel



# Persistent HTTP (HTTP 1.1)

## Persistent HTTP



- Multiple objects can be sent over single TCP connection between client, server
- Server leaves connection open after sending response
- Subsequent HTTP messages between same client/server sent over open connection

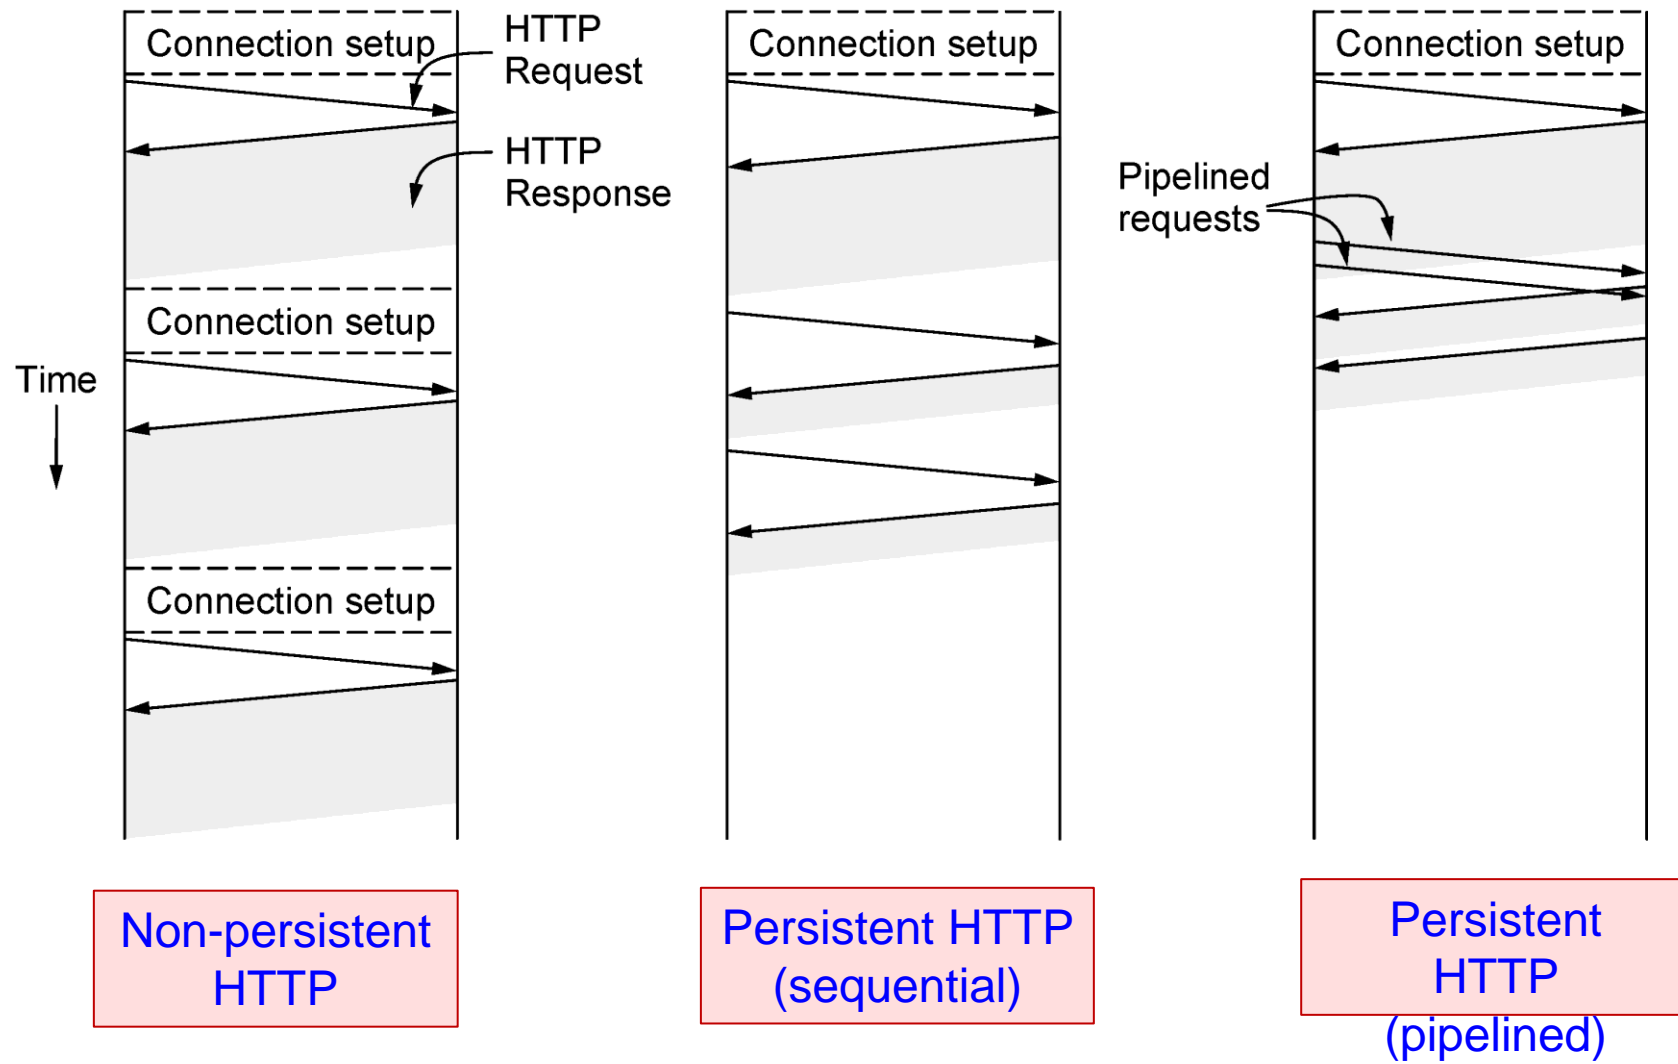


# Persistent HTTP (HTTP 1.1)

## Persistent HTTP

- Non-pipelined:
  - Response time of downloading a HTML file with 5 embedded objects
- Pipelined: Multiple requests can be sent together





# HTTP 1.1

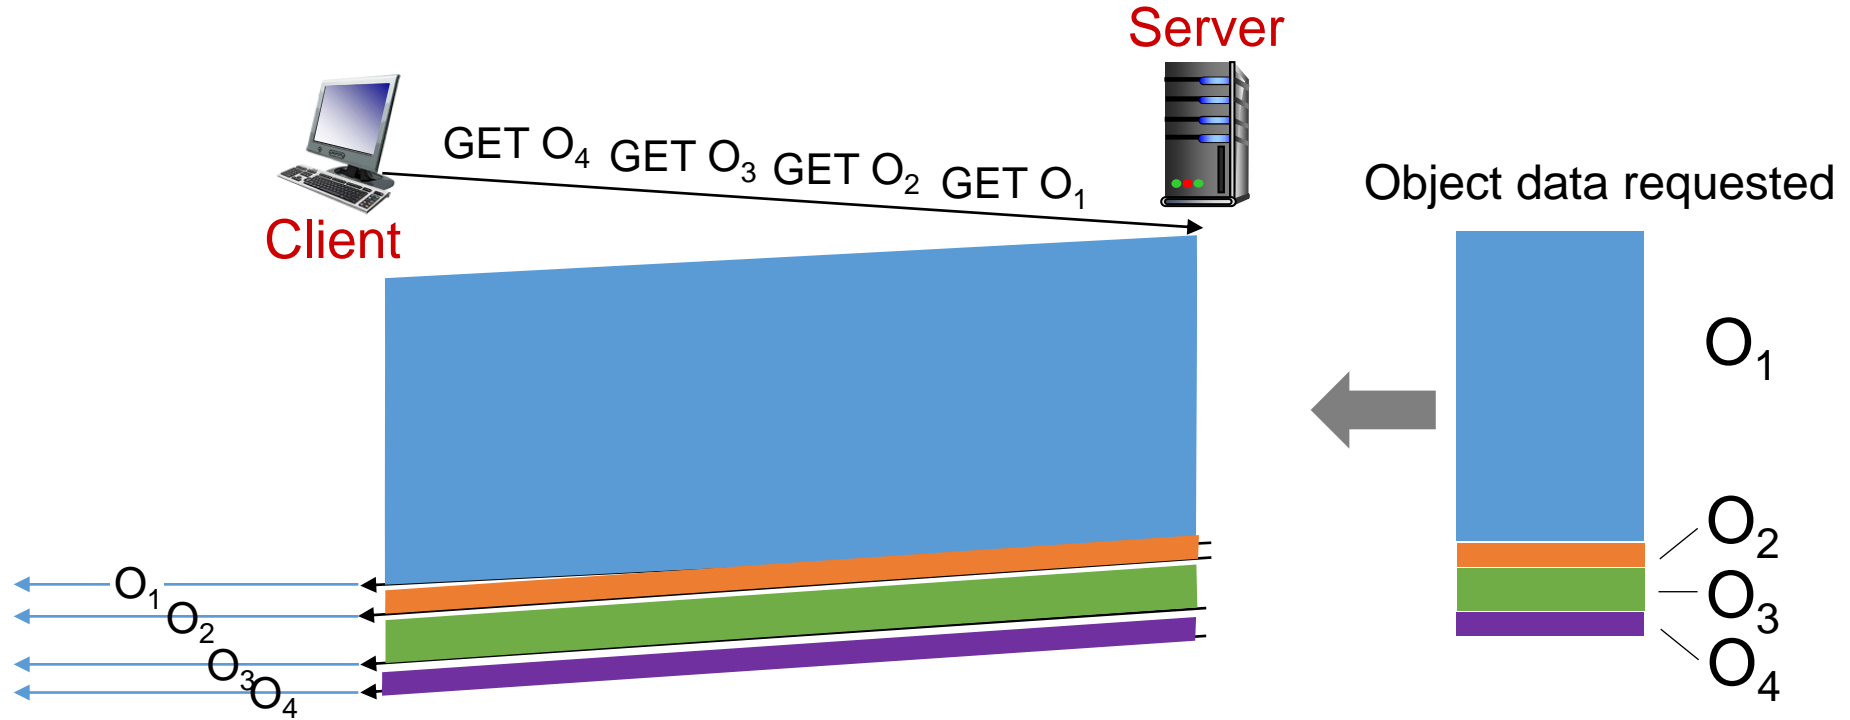
HTTP1.1: Introduced multiple, pipelined GETs over single TCP connection

- Server responds in-order (FCFS: first-come-first-served scheduling) to GET requests
  - With FCFS, small object may have to wait for transmission (head-of-line (HOL) blocking) behind large object(s)
-



# HTTP 1.1 (HOL Blocking)

HTTP 1.1: client requests 1 large object (e.g., video file) and 3 smaller objects



Objects delivered in order requested: `O2`, `O3`, `O4` wait behind `O1`

# HTTP/2

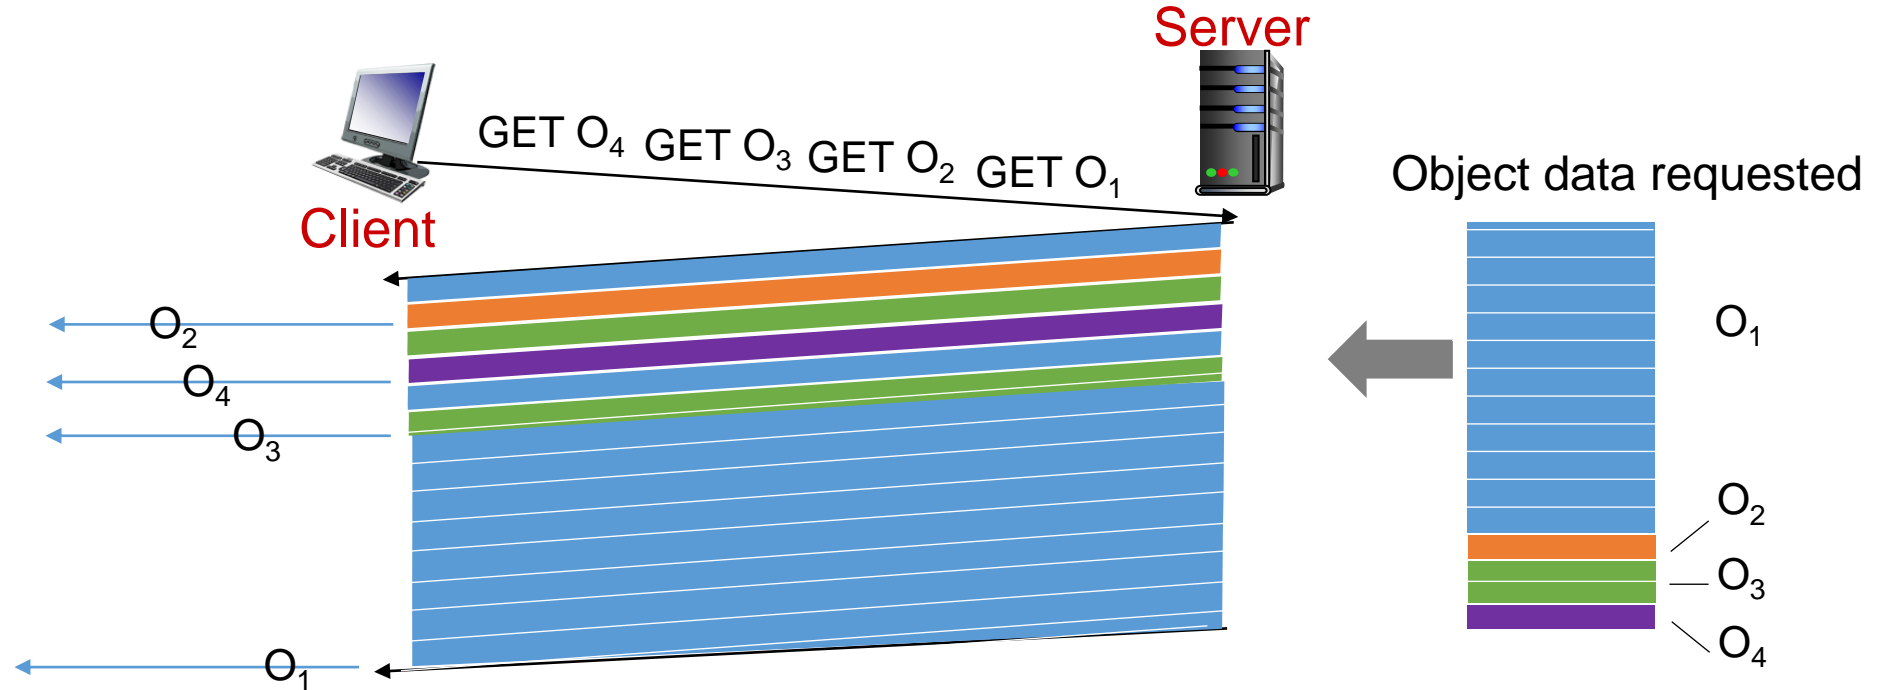
**Key goal:** Decreased delay in multi-object HTTP requests

HTTP/2: [RFC 7540, 2015] Increased flexibility at server in sending objects to client:

- Methods, status codes, most header fields unchanged from HTTP 1.1
  - Transmission order of requested objects based on client-specified object priority (not necessarily FCFS)
  - Push unrequested objects to client
  - Divide objects into frames, schedule frames to mitigate HOL blocking
-

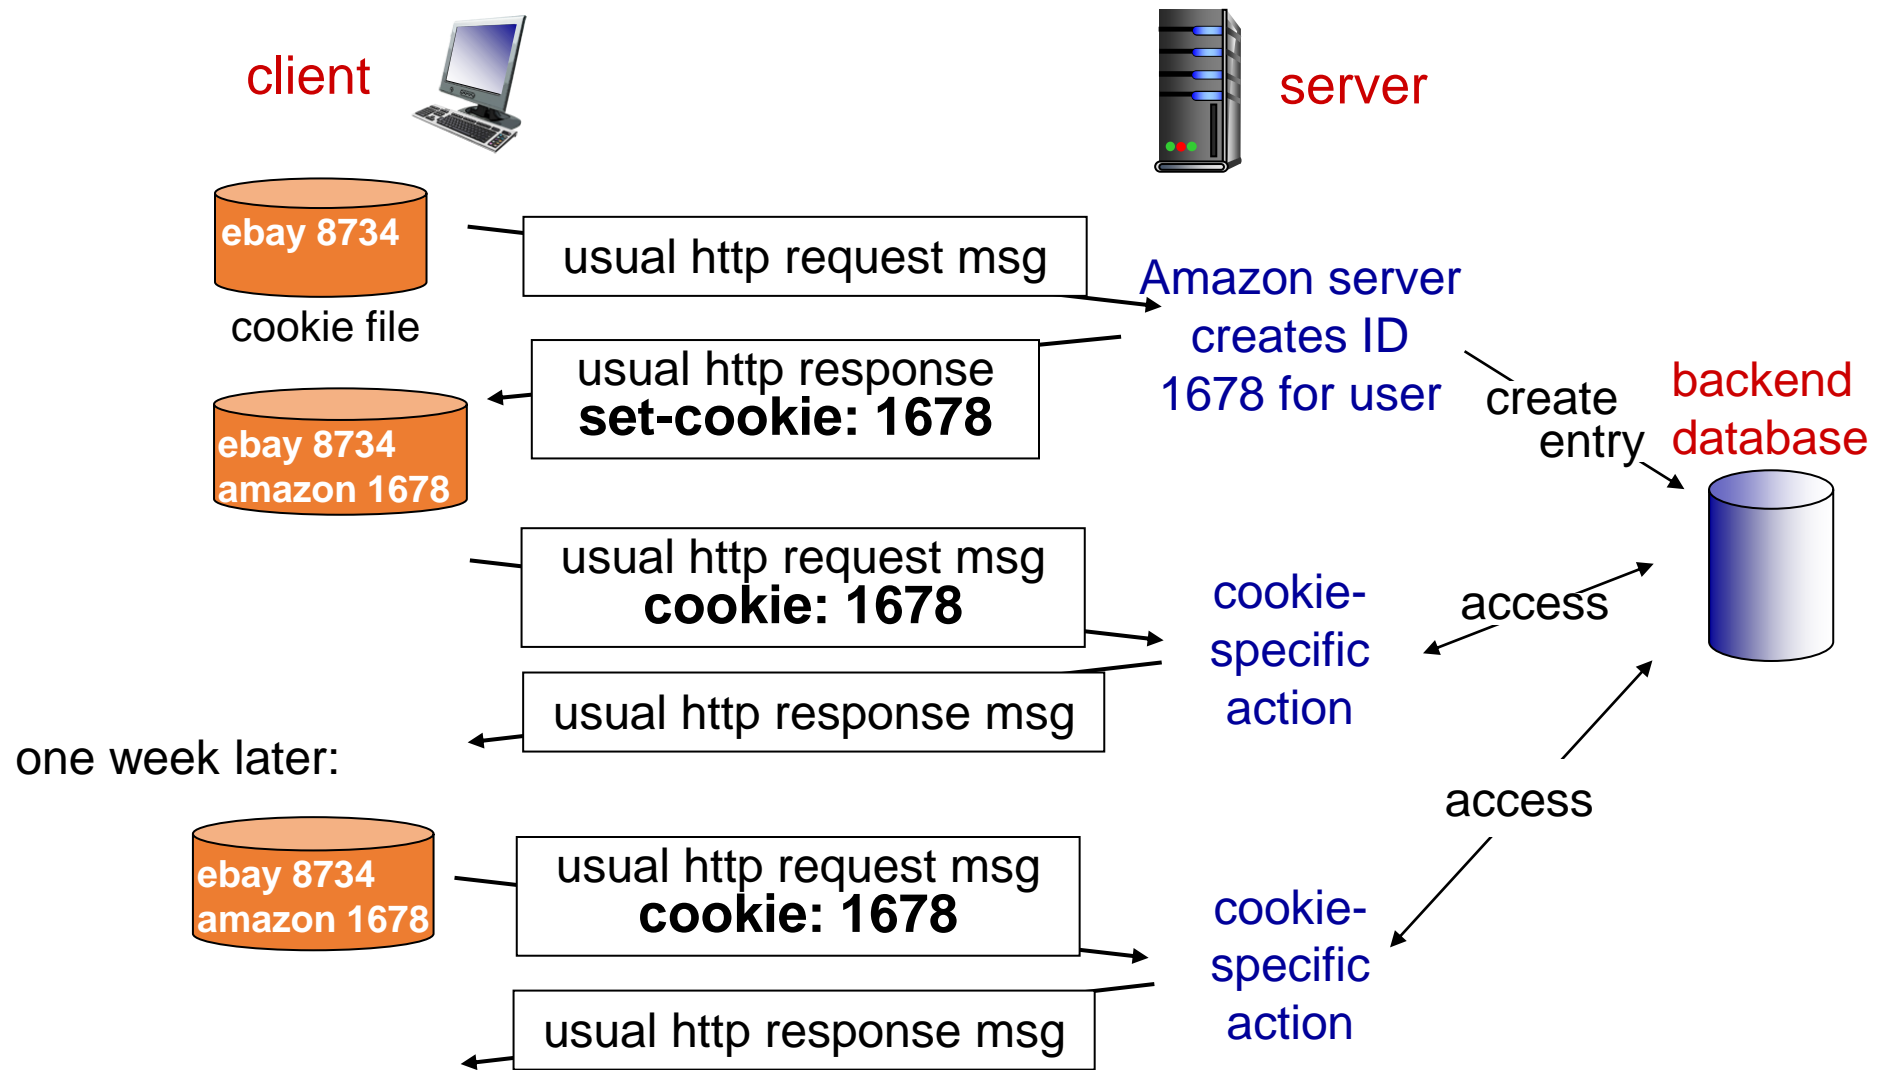
# HTTP/2: Mitigating HOL Blocking

HTTP/2: objects divided into frames, frame transmission interleaved



O<sub>2</sub>, O<sub>3</sub>, O<sub>4</sub> delivered quickly, O<sub>1</sub> slightly delayed

# Cookies



# User-server State: Cookies

Many Web sites use cookies

## Four components:

- 1) Cookie header line of HTTP response message
- 2) Cookie file kept on user's host, managed by user's browser
- 3) Cookie header line in next HTTP request message
- 4) Back-end database at Web site

## Example:

- An user visits specific e-commerce site for first time
  - When initial HTTP requests arrives at site, site creates:
    - Unique ID
    - Entry in backend database for ID
-

# Cookies (continued)

## What cookies can be used for:

- Authorization
  - Shopping carts
  - Recommendations
  - User session state (Web e-mail)
-


# Web Caches

---

# Web caching (Browser cache + Proxy cache)

User often revisit same pages

**Goal:** Satisfy client request without involving origin server






# Web caching (Browser cache + Proxy cache)

## Why Web caching?

- Reduce response time for client request
- Reduce traffic on an institution's access link



# How to verify if cached content is up-to-date?

- Server tells cache about object's allowable caching in response header:

Cache-control: max-age = <seconds>

- Max time until the content is considered fresh

Cache-control: public

- Content is cacheable

Cache-control: private

- Content is cacheable at the browser, but not at the proxy cache

Cache-control: no-cache

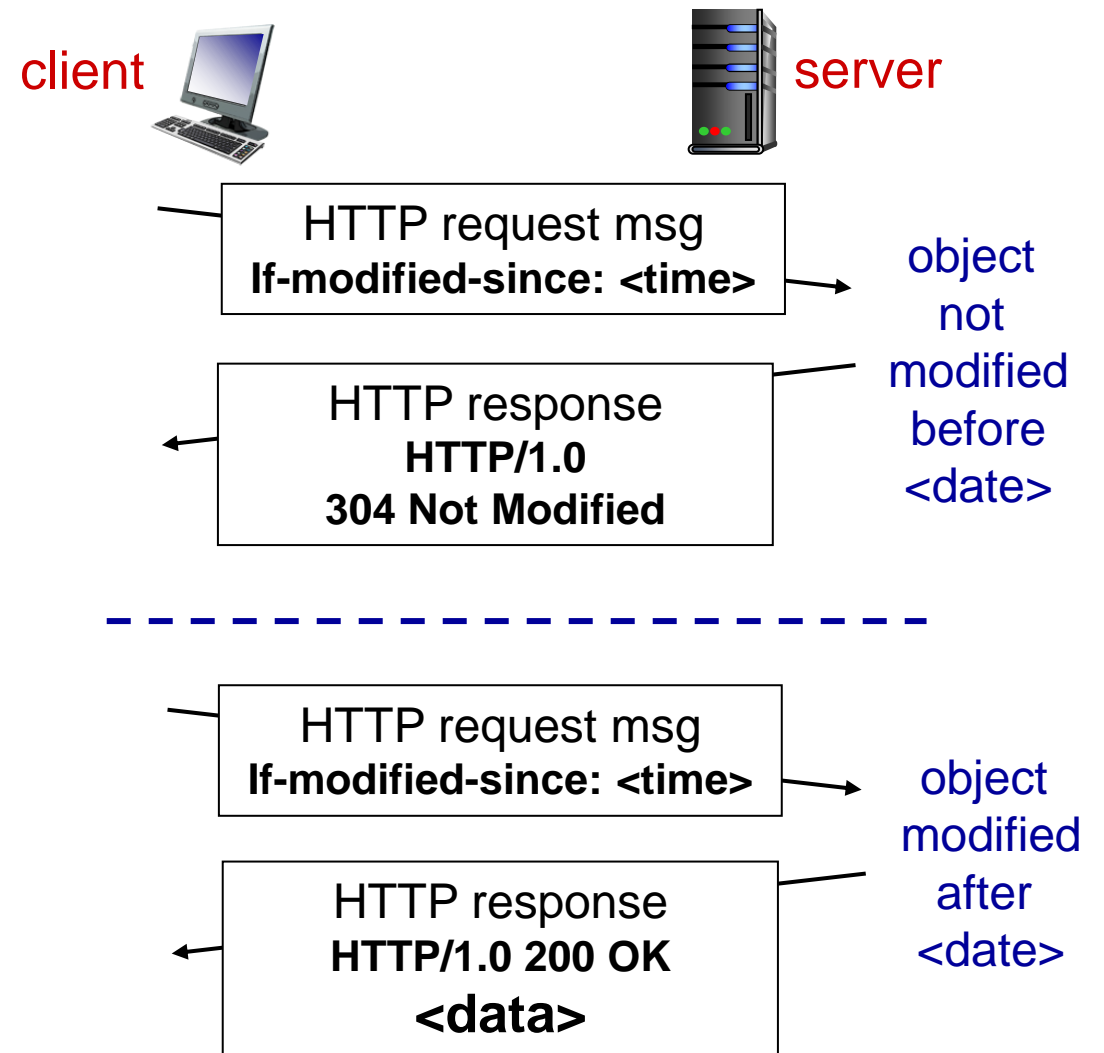
- Cache but revalidate

Cache-control: no-store

- No caching
-

# Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version
- Using **Last-Modified + Conditional GET**
- **Cache:** specify time of cached copy in HTTP request  
**If-modified-since: <time>**
- **Server:** response contains no object if cached copy is up-to-date:  
**HTTP/1.0 304 Not Modified**



# Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version
- Using **ETag + Conditional GET**
- **Cache:** specify Etag (hash) of cached copy in HTTP request  
**If-Node-Match: <ETag>**
- **Server:** response contains no object if cached copy is up-to-date:  
**HTTP/1.0 304 Not Modified**

```
HTTP/1.x 200 OK
Date: Mon, 14 Dec 2009 18:50:22 GMT
Server: Apache
Last-Modified: Sun, 22 Nov 2009 19:58:52 GMT
Etag: "478fb2358f700"
Accept-Ranges: bytes
Vary: Accept-Encoding,User-Agent
Cache-Control: max-age=15552000, public
Content-Length: 1150
Keep-Alive: timeout=6, max=32
Connection: Keep-Alive
Content-Type: image/x-icon
```


```
GET /favicon.ico HTTP/1.1
Host: solariz.de
User-Agent: Mozilla/5.0 (Windows; U; Windows NT
Accept: text/html,application/xhtml+xml,application/javascript;q=0.9,*/*;q=0.8
Accept-Language: de-de,de;q=0.8,en-us;q=0.5,en;q=0.3
Accept-Encoding: gzip,deflate
Accept-Charset: ISO-8859-1,utf-8;q=0.7,*;q=0.7
Keep-Alive: 300
Connection: keep-alive
Cookie: wp-settings-2=hidetb%3D1%26editor%3D1
If-Modified-Since: Sun, 22 Nov 2009 19:58:52 GMT
If-None-Match: "478fb2358f700"
```

```
HTTP/1.x 304 Not Modified
Date: Mon, 14 Dec 2009 18:50:32 GMT
Server: Apache
Connection: Keep-Alive
Keep-Alive: timeout=6, max=32
Etag: "478fb2358f700"
Cache-Control: max-age=15552000, public
Vary: Accept-Encoding,User-Agent
```

Src: [https://commons.wikimedia.org/wiki/File:Etag\\_header\\_beispiele.png](https://commons.wikimedia.org/wiki/File:Etag_header_beispiele.png)

# Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version




# Web caching (Web Proxies)

## Scenario:

- Access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- Web object size: 100K bits
- Average request rate from browsers to origin servers: 15/sec
  - Avg data rate to browsers: 1.50 Mbps

## Performance:


- Access link utilization = .97
- LAN utilization: .0015
- End-end delay = RTT + access link delay + LAN delay  
= 2 sec + minutes + usecs



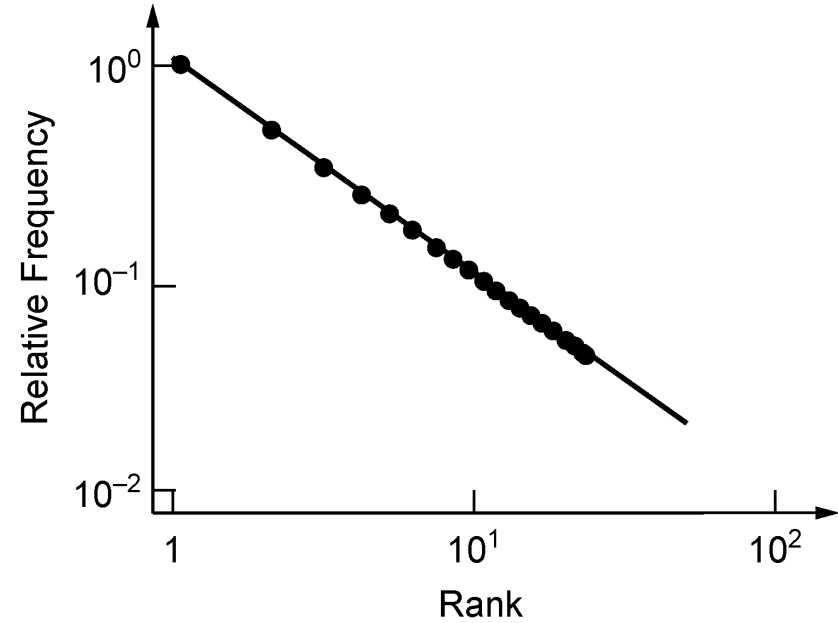
# Internet Content Popularity

$$f(i, \alpha, \mathcal{M}) = \frac{\frac{1}{i^\alpha}}{\sum_{j=1}^{\mathcal{M}} \frac{1}{j^\alpha}} = \frac{1}{\mathbf{H}_{\mathcal{M}, \alpha}}$$

$$\log f(i, \alpha, \mathcal{M}) = \log \left( \frac{1}{\mathbf{H}_{\mathcal{M}, \alpha}} \right) - \alpha \log i$$



Zipf distribution (linear scale)




Zipf distribution (log-log scale)

# Web caching (Web Proxies)

Suppose cache hit rate is 0.4:

- 40% requests served by cache, with low (msec) delay
- 60% requests satisfied at origin
  - Rate to browsers over access link  
 $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
  - Access link utilization  $= 0.9 / 1.54 = .58$  means low (msec) queueing delay at access link
- Average end-end delay:  
 $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache}) = 0.6 (2 \text{ secs}) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$





# Content Distribution Networks (CDNs)


---

# Video Streaming and CDNs

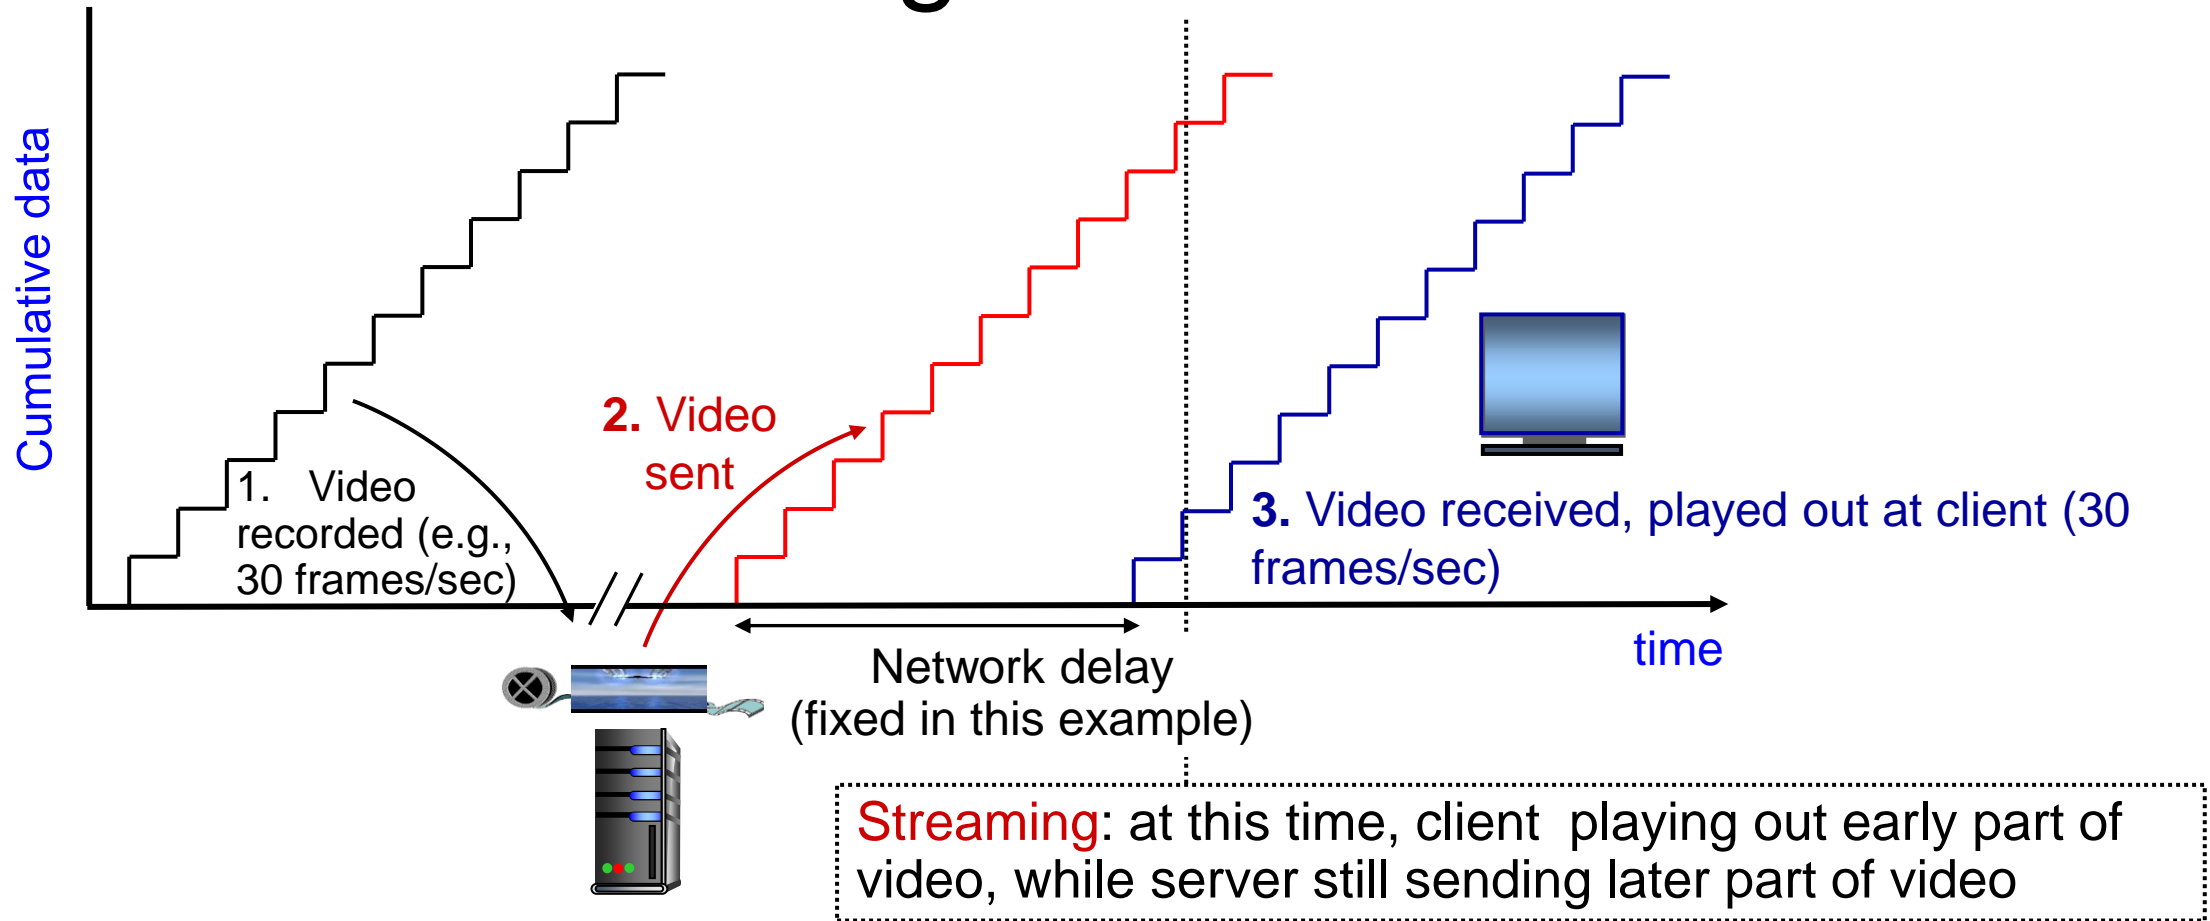
- Stream video traffic: major consumer of Internet bandwidth
  - Netflix, YouTube, Amazon Prime: 80% of residential ISP traffic (2020)
- Challenges:
  - Scale - how to reach ~1B users?
  - Heterogeneity
    - Different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)



# Streaming Stored Video




# Streaming Stored Video




# Streaming Stored Video: Challenges

- **Continuous playout constraint:** During client video playout, playout timing must match original timing
  - But **network delays are variable** (jitter), so will need **client-side buffer** to match continuous playout constraint
- Other challenges:
  - Client interactivity: pause, fast-forward, rewind, jump through video
  - Video packets may be lost, retransmitted



# Streaming Stored Video: Playout Buffering



- **Client-side buffering and playout delay:** compensate for network-added delay, delay jitter

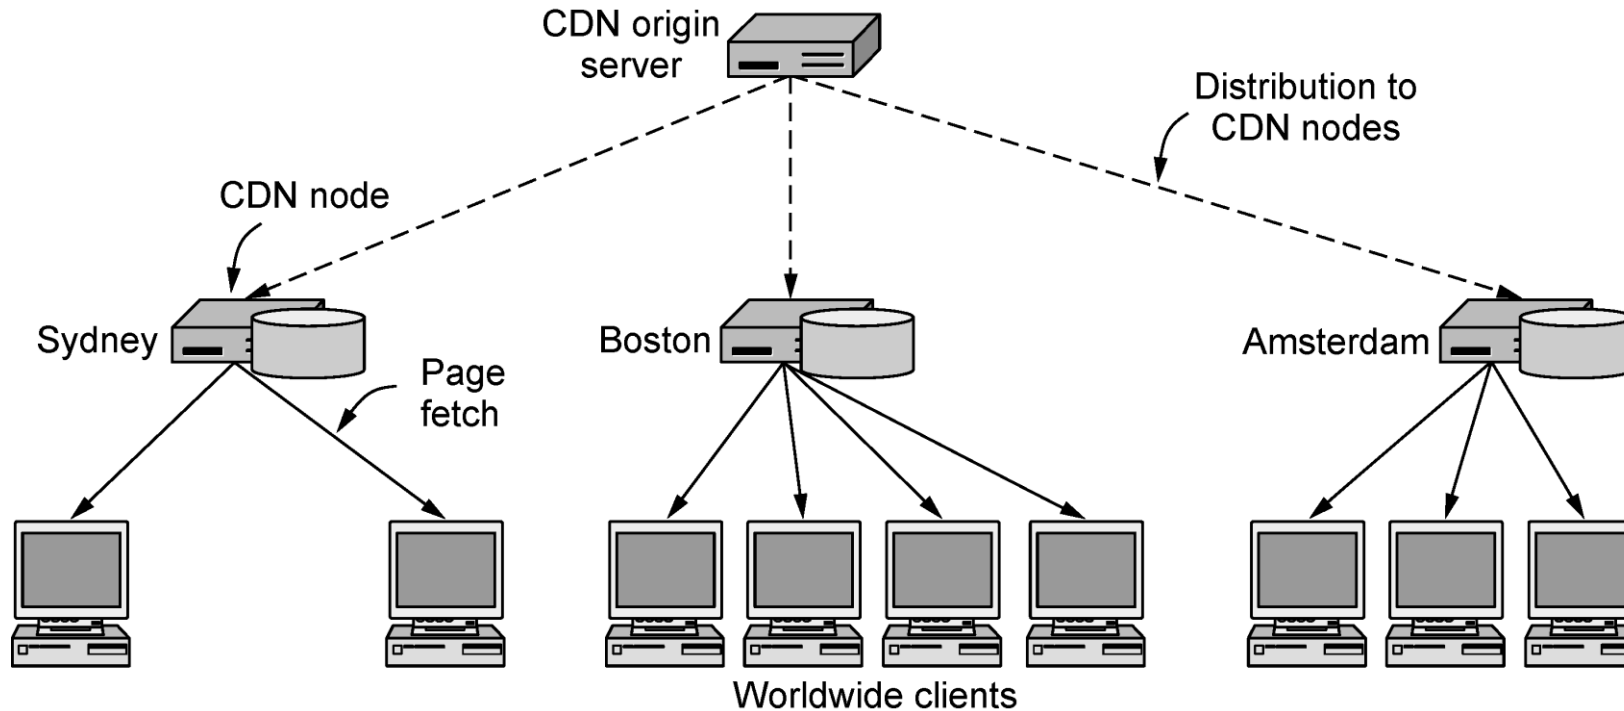
# Content Distribution Networks (CDNs)

**Challenge:** How to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?

- **Solution:** Store/serve multiple copies of videos at multiple geographically distributed sites (**CDN**)
- Push CDN servers deep into many access networks
  - Close to users
  - Akamai: 240,000 servers deployed in > 120 countries (2015)

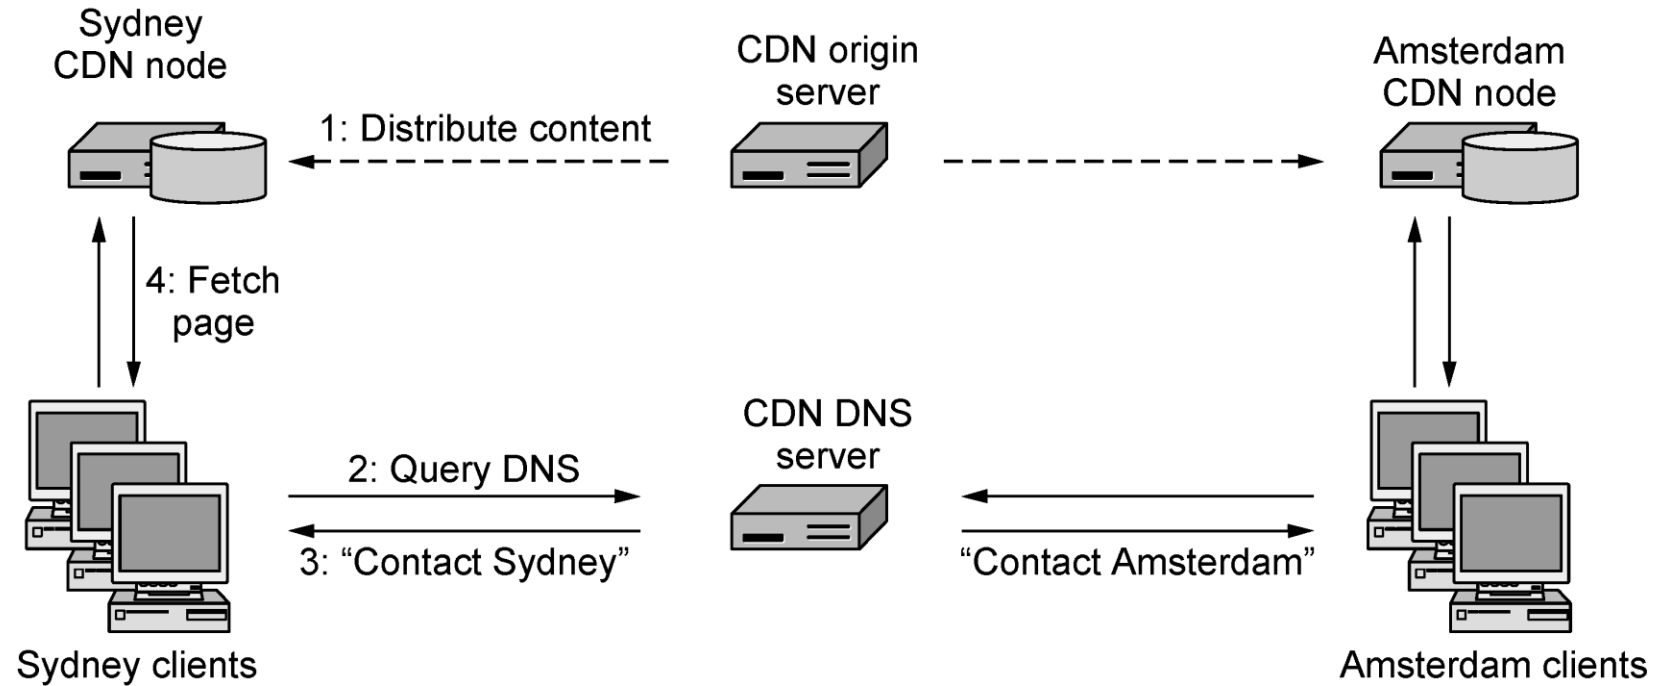


# Content Distribution Networks (CDNs)





# Content Distribution Networks: DNS Redirection



<http://www.Netflix.com/page.html>

# Content Distribution Networks (CDNs)

Win-win solution for all parties:

- Reduces server load → Content providers are 😊
  - Reduces network load → ISPs are 😊
  - Reduces content download time → Users/customers are 😊
-

# Electronic Mail

---


# Electronic Mail

## User Agent

- a.k.a. “mail reader”
- Composing, editing, reading mail messages

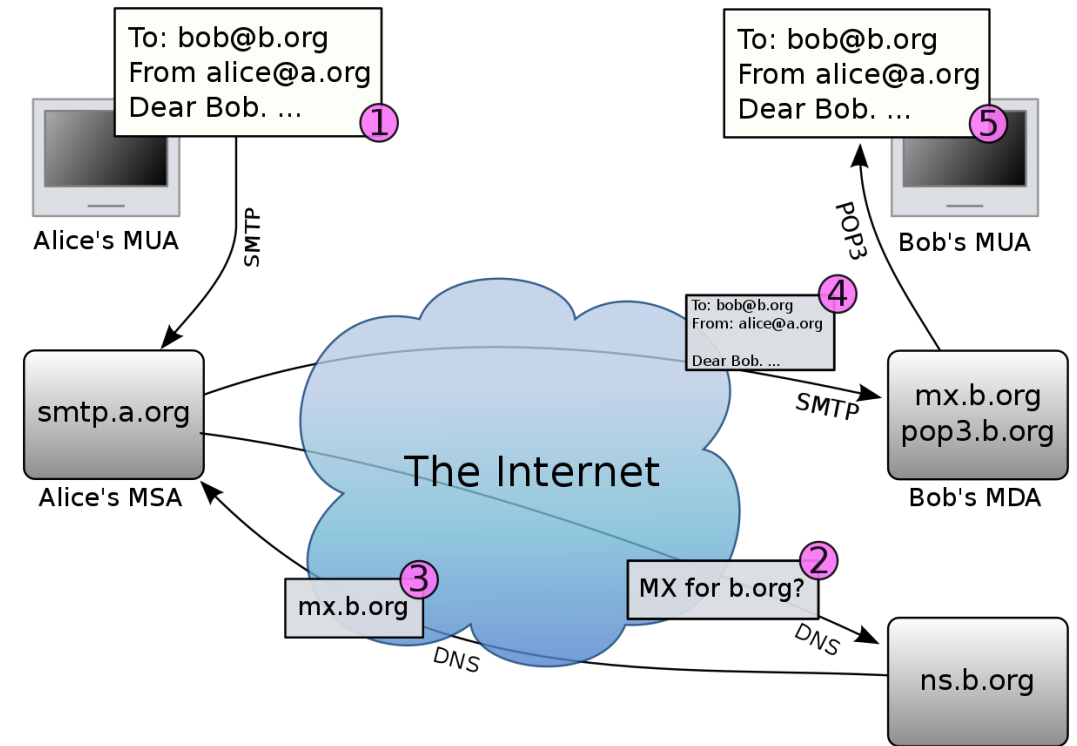
## Mail servers:

- Mailbox contains incoming messages for user
- Message queue of outgoing (to be sent) mail messages



# Electronic Mail

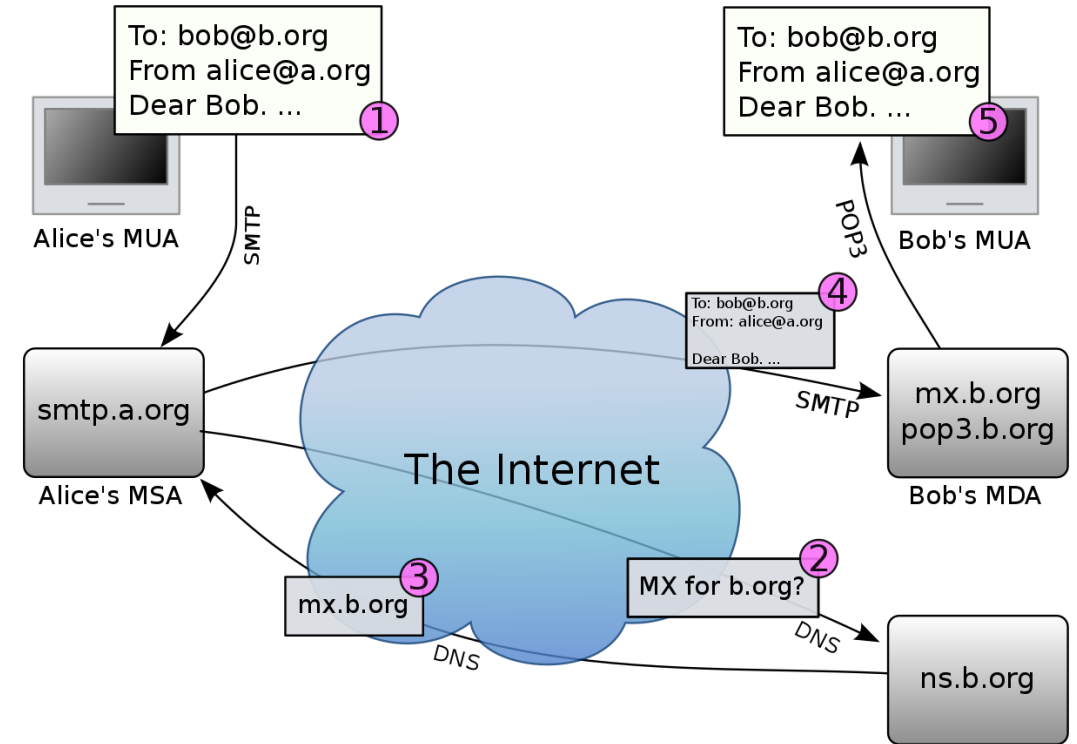
- **Mail transfer protocol**: Delivery/storage to receiver's server → **Push** operation
  - **SMTP**: Simple Mail Transfer Protocol
    - User agent to mail server
    - In between mail servers
- **SMTP protocol** between mail servers to send email messages
  - Client: sending mail server
  - “Server”: receiving mail server



Src: [https://en.wikipedia.org/wiki/Open\\_mail\\_relay#/media/File:Email.svg](https://en.wikipedia.org/wiki/Open_mail_relay#/media/File:Email.svg)

# Electronic Mail

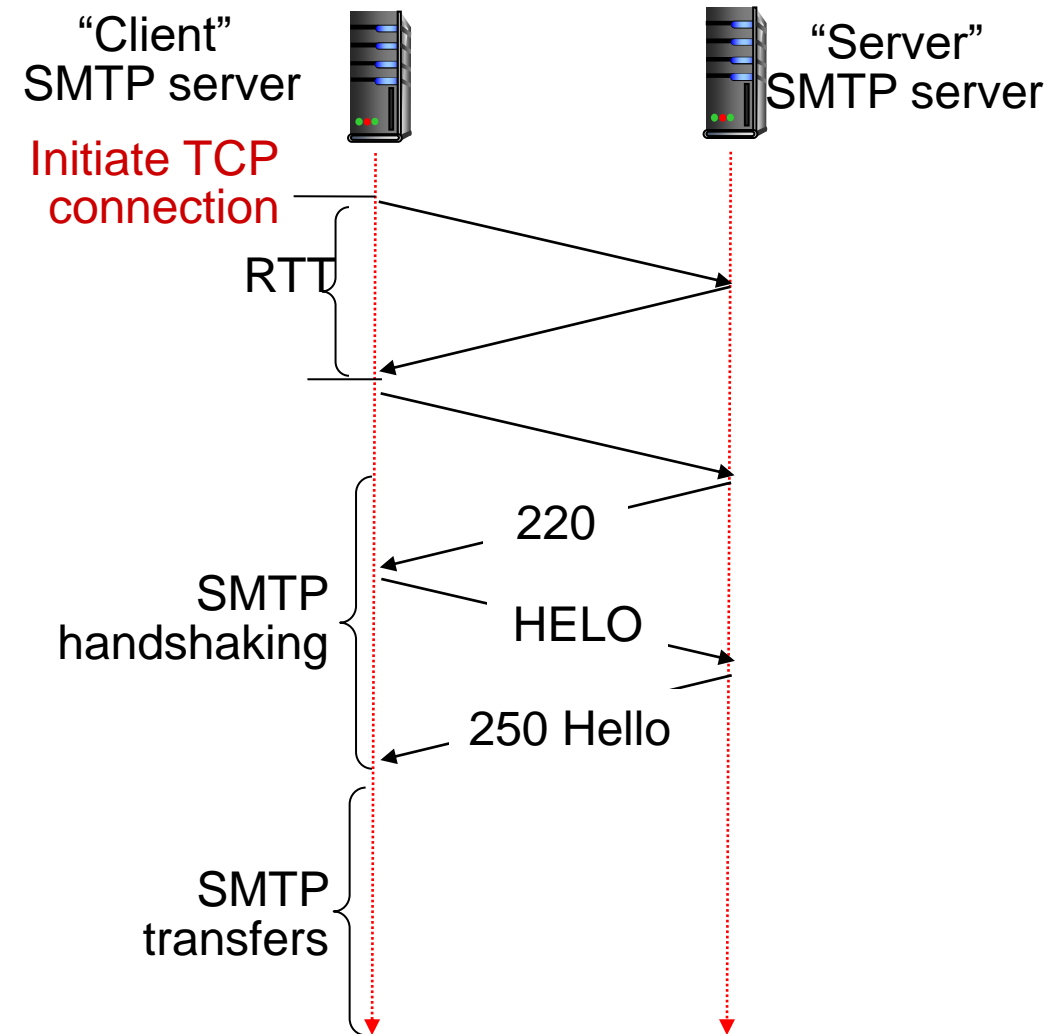
- **Mail access protocol:** retrieval from server → **Pull** operation
  - **POP:** Post Office Protocol [RFC 1939]: authorization, download
  - **IMAP:** Internet Mail Access Protocol [RFC 1730]: more features, including manipulation of stored messages on server
  - **HTTP:** gmail, Hotmail, Yahoo! Mail, etc.



Src: [https://en.wikipedia.org/wiki/Open\\_mail\\_relay#/media/File:Email.svg](https://en.wikipedia.org/wiki/Open_mail_relay#/media/File:Email.svg)

# SMTP

- Uses TCP to reliably transfer email message from client (mail server initiating connection) to server, port 25
  - Direct transfer: sending server (acting like client) to receiving server
- Three phases of transfer
  - SMTP handshaking (greeting)
  - SMTP transfer of messages
  - SMTP closure



# Summary

## □ Web and HTTP:

- Persistent and non-persistent HTTP
- Web cache

## □ Electronic mail:

- SMTP, POP, IMAP
-